

Spark & Cassandra

Jose María Muñoz Rey

DevFest Córdoba
2016

Who am I?



- Software Engineering Student.
- Data analytics and processing @ s|ngular.
- Community enthusiast.

s|ngular

 GDG Cáceres



The Spark logo features the word "Spark" in a white, sans-serif font, with a white star icon positioned above the letter "k".


Spark

The Cassandra logo consists of a white, stylized eye icon with long, radiating eyelashes, positioned above the word "cassandra" in a white, lowercase, sans-serif font.

cassandra

About talk?

- Motivation
- Cassandra
- Spark
- Spark + Cassandra
- Demo
- Conclusions



The need to lead in data and analytics

What do you need?

- Decent analytical query performance.
- Big throughput of reads/writes.
- New ways of do data analytics.
- Strong consistency is not necessary.
- But Big availability!
- Combine more than one data source.



+



Apache Cassandra

- Google BigTable + Amazon DynamoDB.
- No single point of failure (Masterless).
- “Linear” scalability.
- Support multi-data center.
- CAP Theorem - AP.
- Consistency level on query level.
- Maintained by Datastax.



Composition

Cluster

Datacenter A

Rack A

Nodo 8

Nodo 9

Nodo 0

Datacenter B

Rack B

Nodo 1

Nodo 2

Nodo 3

Nodo 4

Rack C

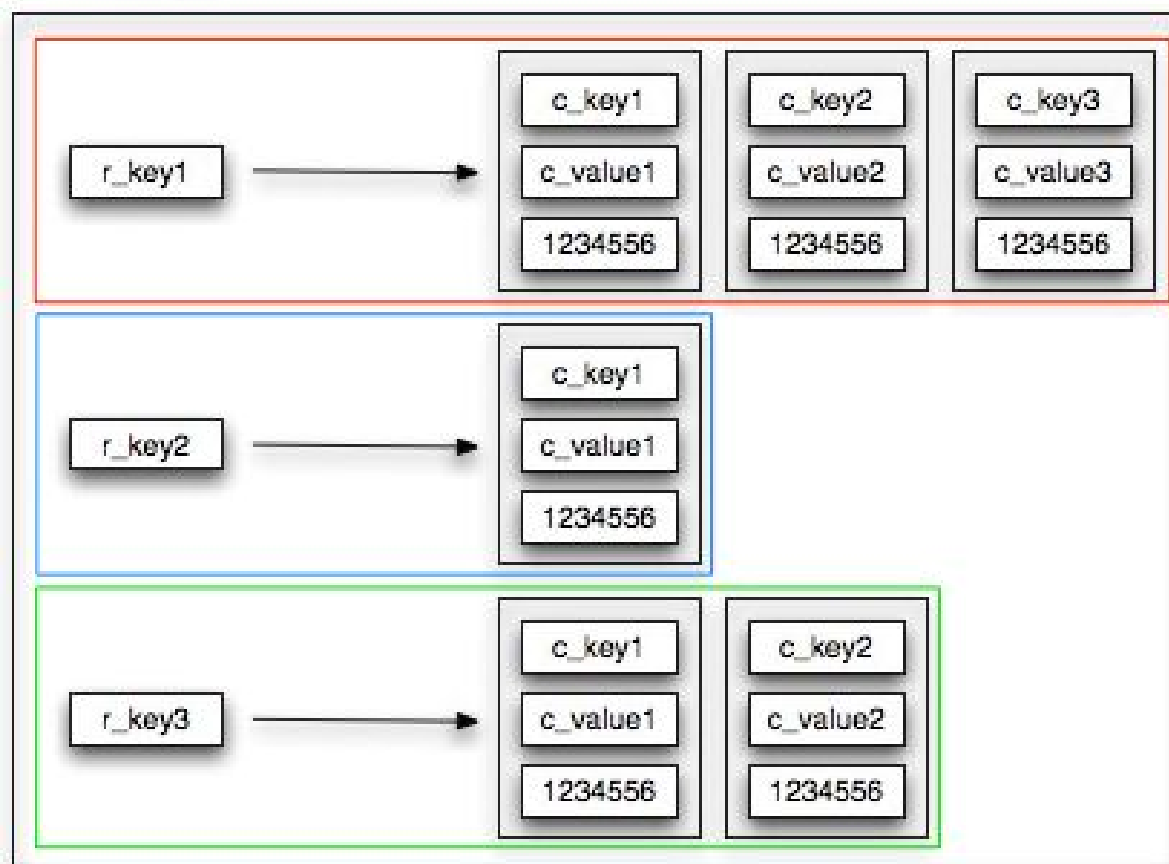
Nodo 5

Nodo 6


Nodo 7

...

Data Modelling



About data



- PRIMARY KEY = PARTITION KEY + CLUSTERING KEY
 - PARTITION KEY (hashed)
 - Defines the data locality.
 - Can be compound.
 - CLUSTERING KEY
 - Define the order inside the partition
- Any node can coordinate a query.
- The driver usually have the partition function.

Be careful ***

CQL

Basically like standard SQL but significantly limited.

- Restrictions to avoid full-table queries. ***
- We must know how the data is stored.

But it has some nice functionalities.

- **Complex data types:** JSON, Collections (set, map...) and user defined.
- **User defined functions*** in Java/JavaScript.
- **Secondary index** (columns, any value, even collections)***.
- **Materialized views.**

Querying data

Amount of data queried:

- Or full partition key (=, IN).
- Or Indexes (limited).

Sorting:

- Remember the clustering column? Use it! (<,>,>=,<=)

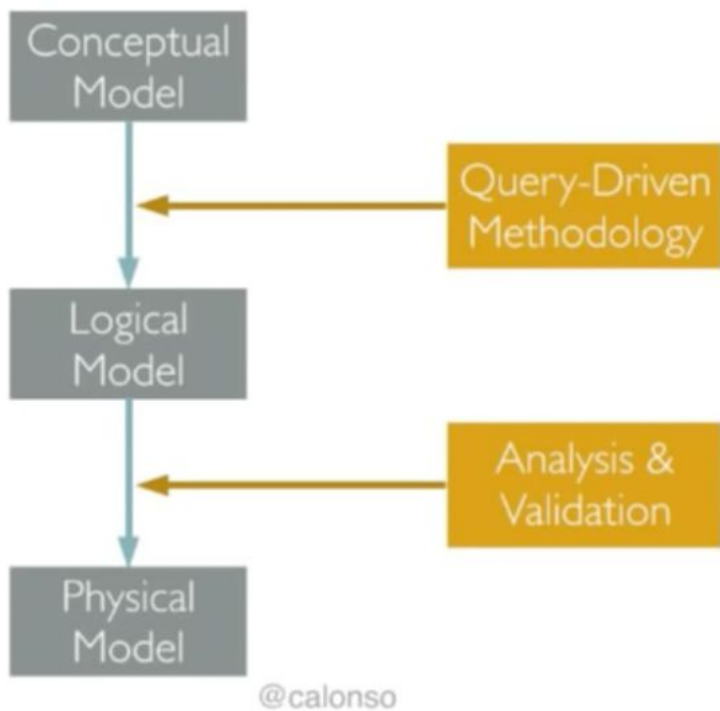
Other:

- JSON, CSV output/input support.
- No JOIN, GROUP BY.

((id, day), time), measure

```
SELECT *  
FROM iot  
WHERE id = "159323"  
AND day = "2016-12-03"  
AND "14:00:00.000" < time  
AND time < "20:00:00.000";
```

Query-Driven Methodology



- Start from Entity-relationship diagram.
- Query-Driven Methodology.
- Logical model (Chebotko).

cyclist_category	
category	K
id	Cl
points	
lastname	

← partition key
← clustering column

- Analysis & Validations.
- Physical implementation.
- **The Kashlev Data Modeler.**

Replication

- **Replication Factor defined on keyspace:**

```
CREATE KEYSPACE my_database
```

```
WITH REPLICATION = {
```

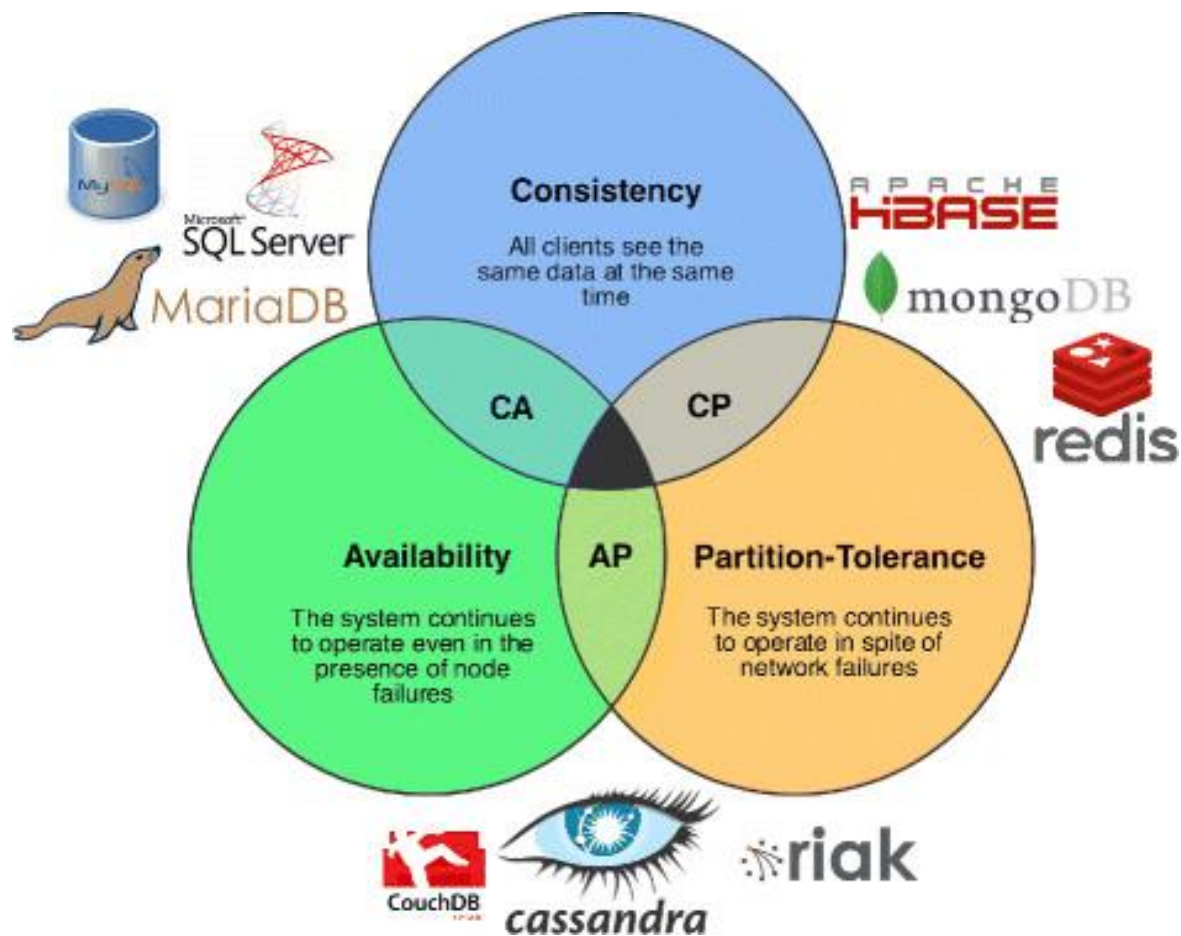
```
    "Class" : "networkTopologyStrategy",
```

```
    "Dc-east" : 2, "dc-west" : 3
```

```
};
```

- Multiple class styles to support multiple schemas.
- **Data close to the users.**
- **Functionality segregation** (DC for users, DC for analytics).

CAP Theorem



Consistency



Availability

Consistency VS Availability

- **Real State vs Velocity**

On **Keyspace** Creation


- Replication factor

On **Query**

- ConsistencyLevel - ONE, QUORUM “ $(Total/2)+1$ ”, ALL


Lightweight Transactions

- Compare-and-set operations




Scale-Up Linearity

Client Writes/s by node count – Replication Factor =



NETFLIX



Widely used

Company	Nodes	Data Volume	Operations
Apple	75.000	10 PB	
Netflix	2.500	420 TB	1 billion/day
Easou	270	300 TB	800 millions/day
eBay	100	250 TB	500 millions/day

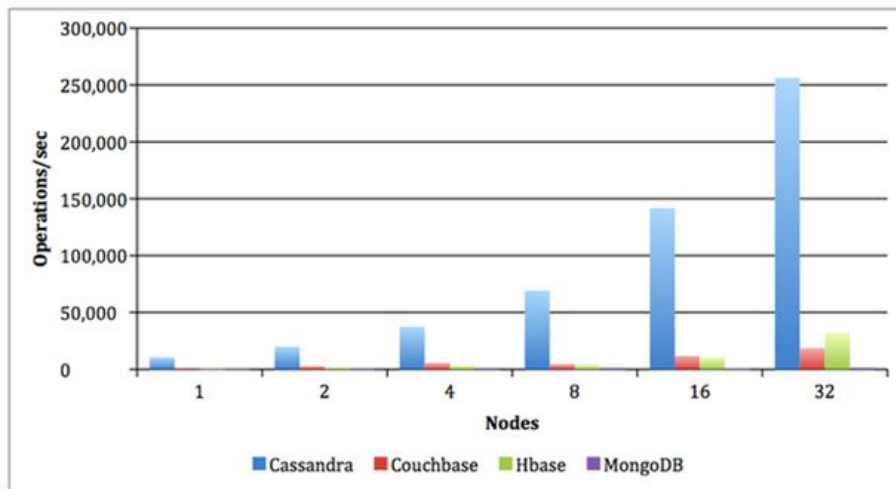


NETFLIX

eaSou 宜搜

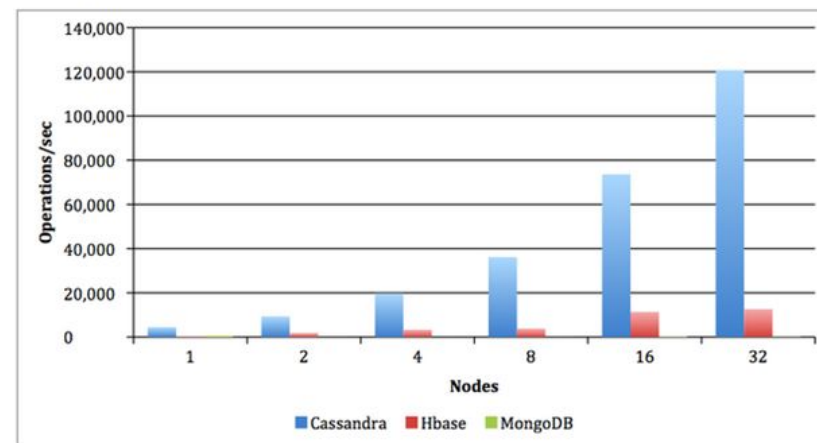
ebay

Read-Modify-Write Workload

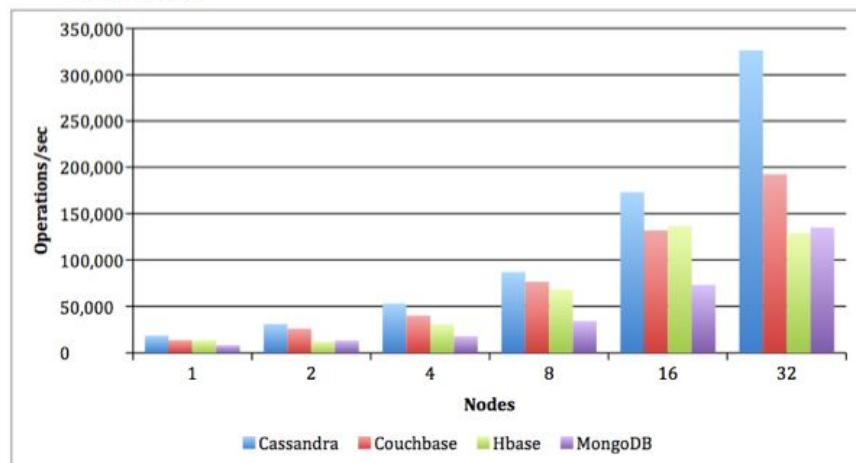


Mixed Operational and Analytical Workload


Note that Couchbase was eliminated from this test because it does not support scan operations (producing the error: "Range scan is not supported").



Load Process



Read-mostly Workload



Conclusions

The Bads

- Remember Query Driven Methodology, be careful.
- Queries have low flexibility.
- Use secondary Indexes usually are a bad practice.

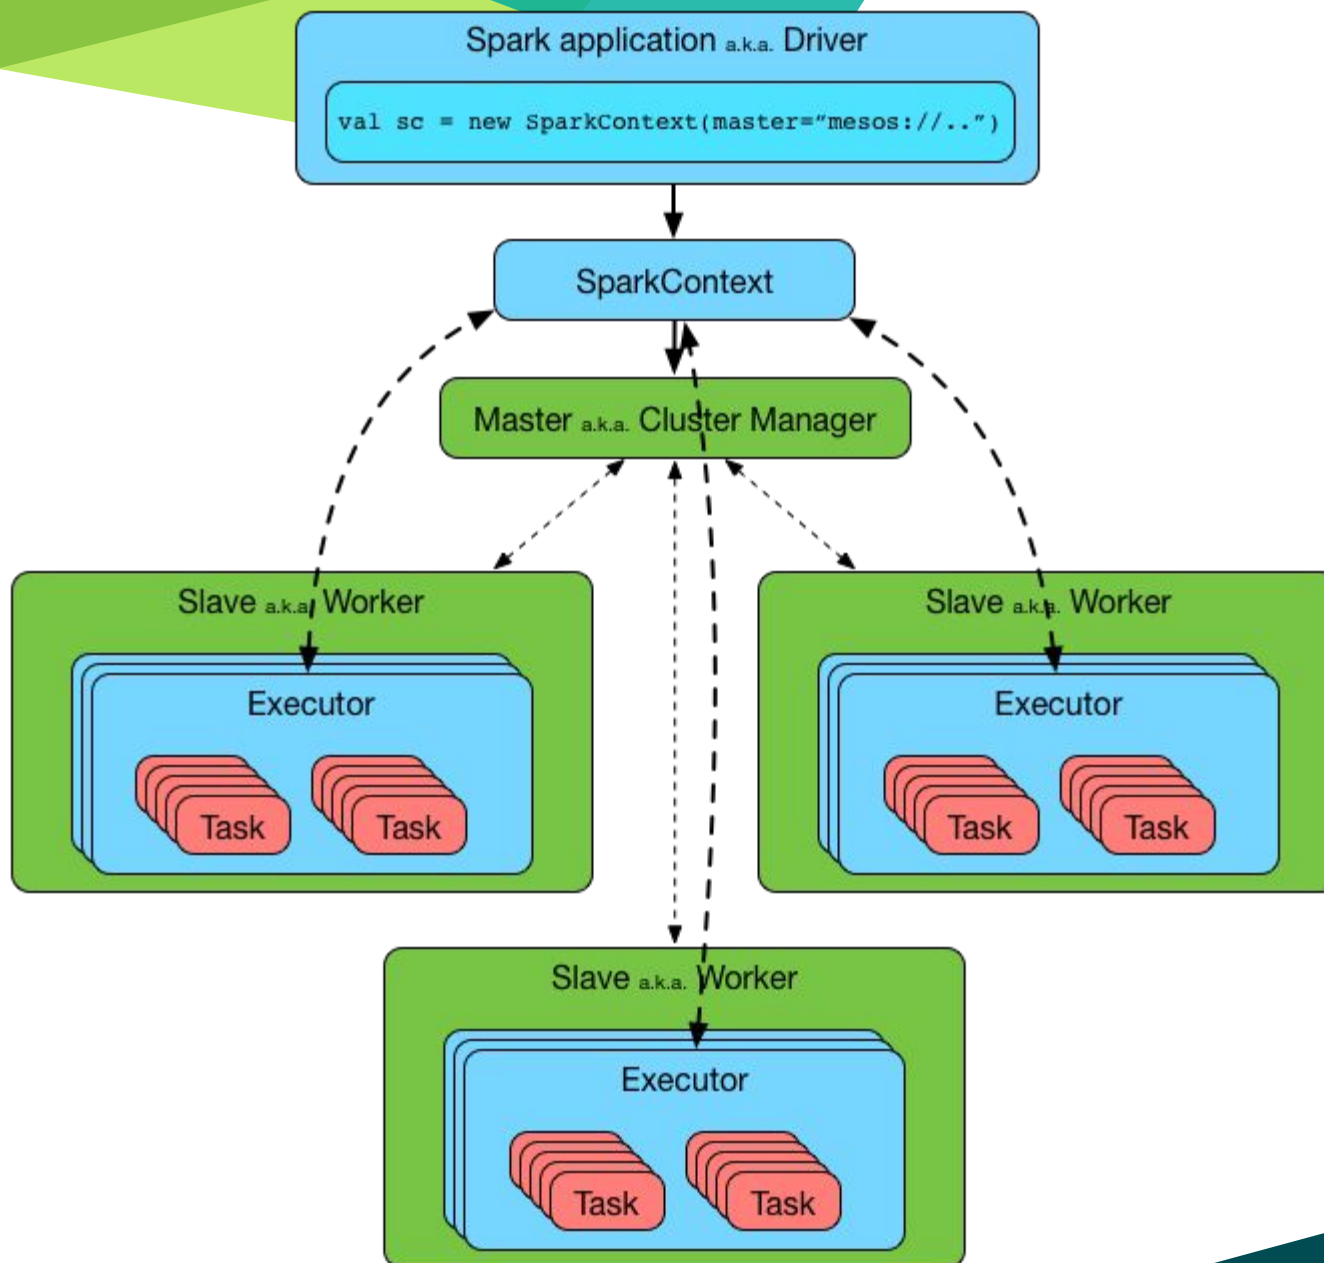
The Goods

- Linear scaling.
- Storage is cheaper than processing.



Apache Spark

- Framework for distributed data processing.
- Fault tolerant.
- Faster than Hadoop (10x-100x).
- Many integrations with frameworks, libraries, databases...
- Batches and streaming analytics.
- Machine Learning with distributed processing.
- Easy to use.



Spark
SQL

Spark
Streaming

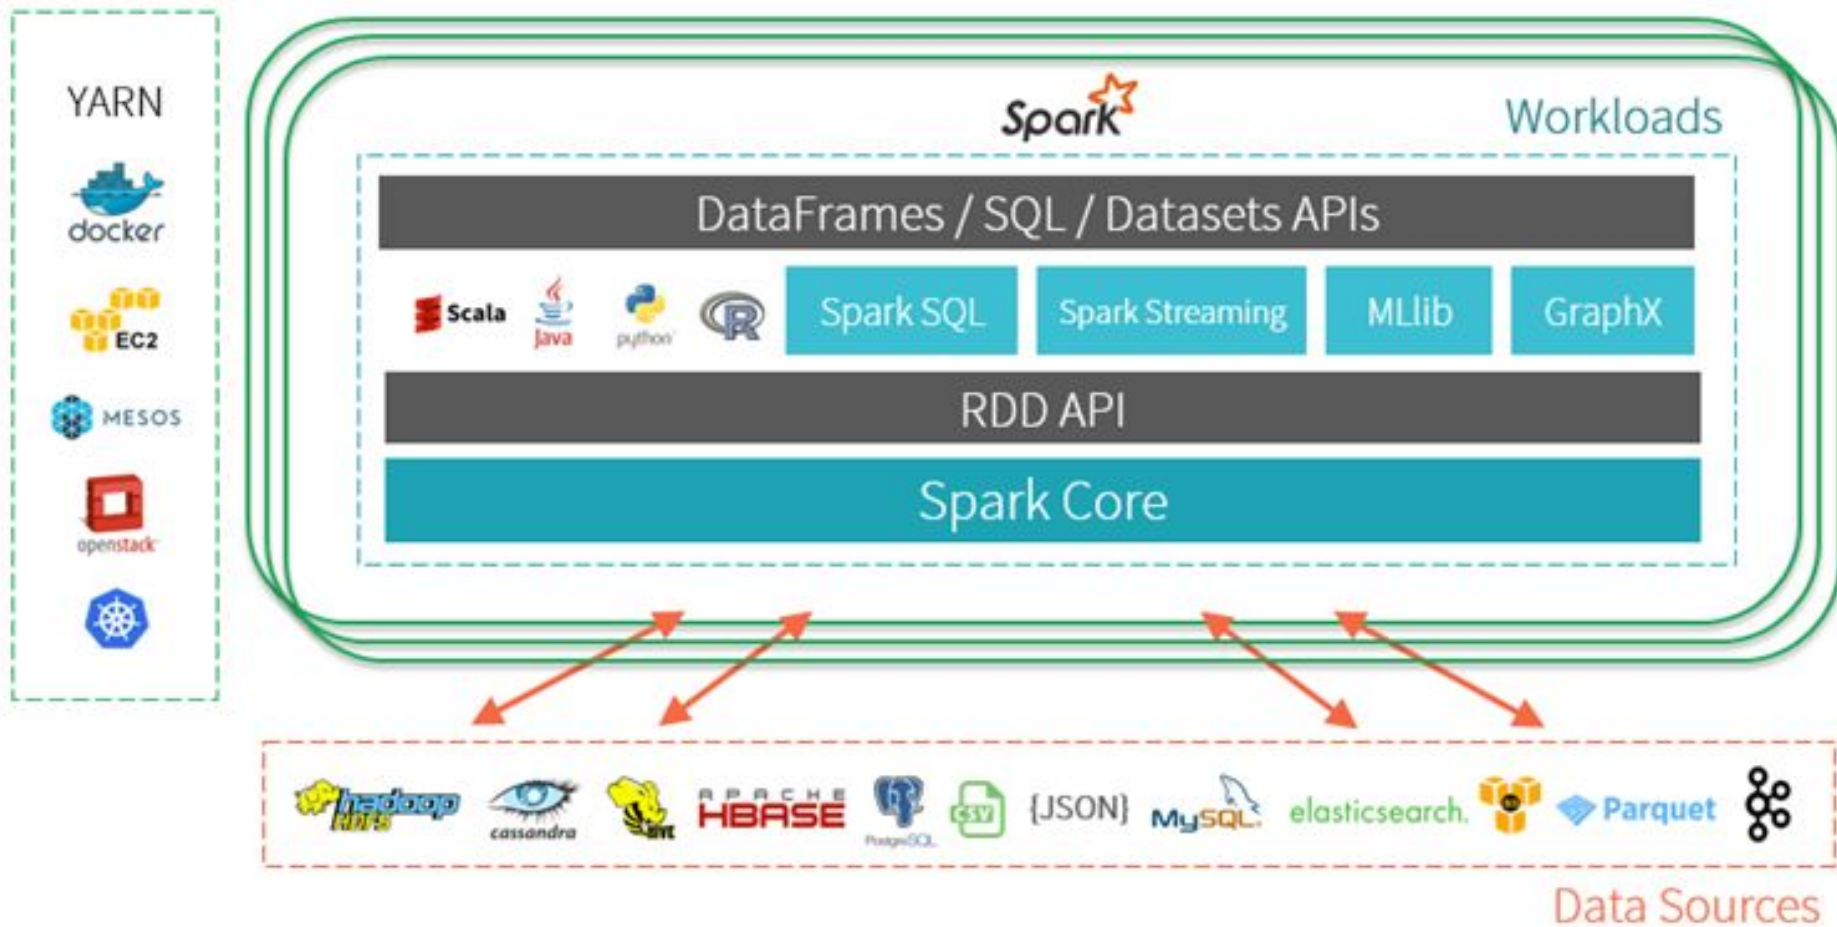
MLlib
(machine
learning)

GraphX
(graph)


Apache Spark

Goal: unified engine across data **sources**,
workloads and environments

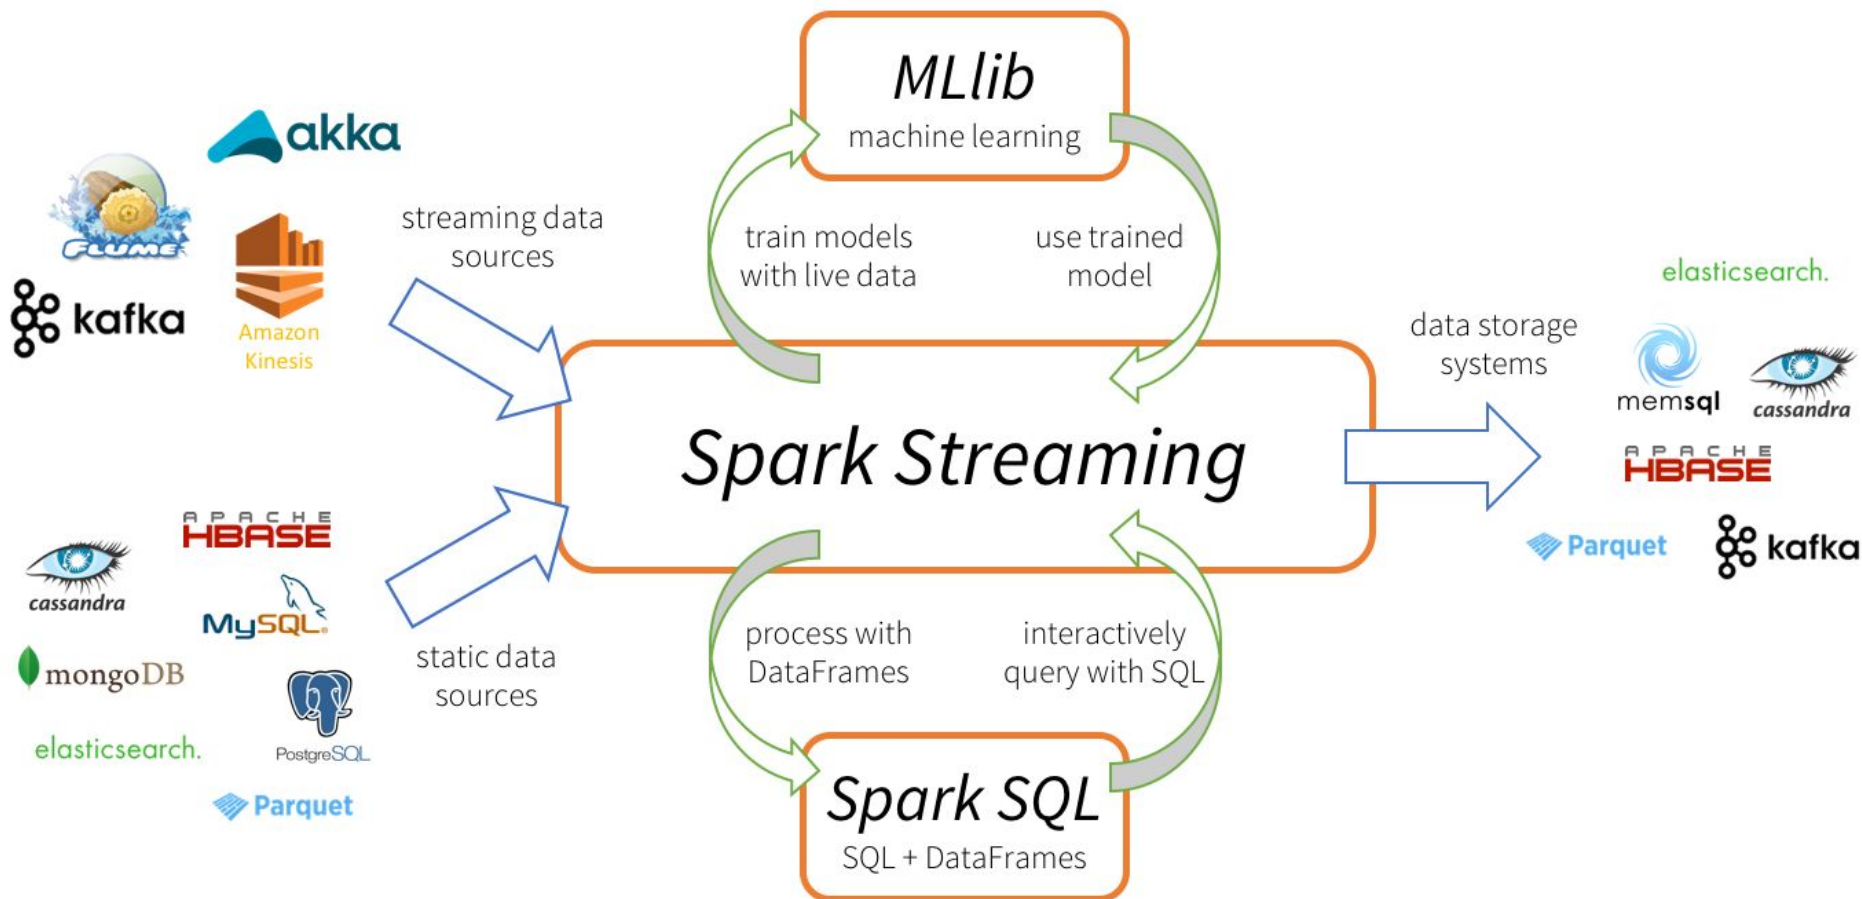
Environments



RDD, DataFrame & Dataset



- RDD: Resilient Distributed Dataset.
- DataFrame = RDD + Schema (SQL!).
- Dataset = DataFrame + Language.
 - Type-safe.
 - Object-oriented programming interface.



SQL or DSL

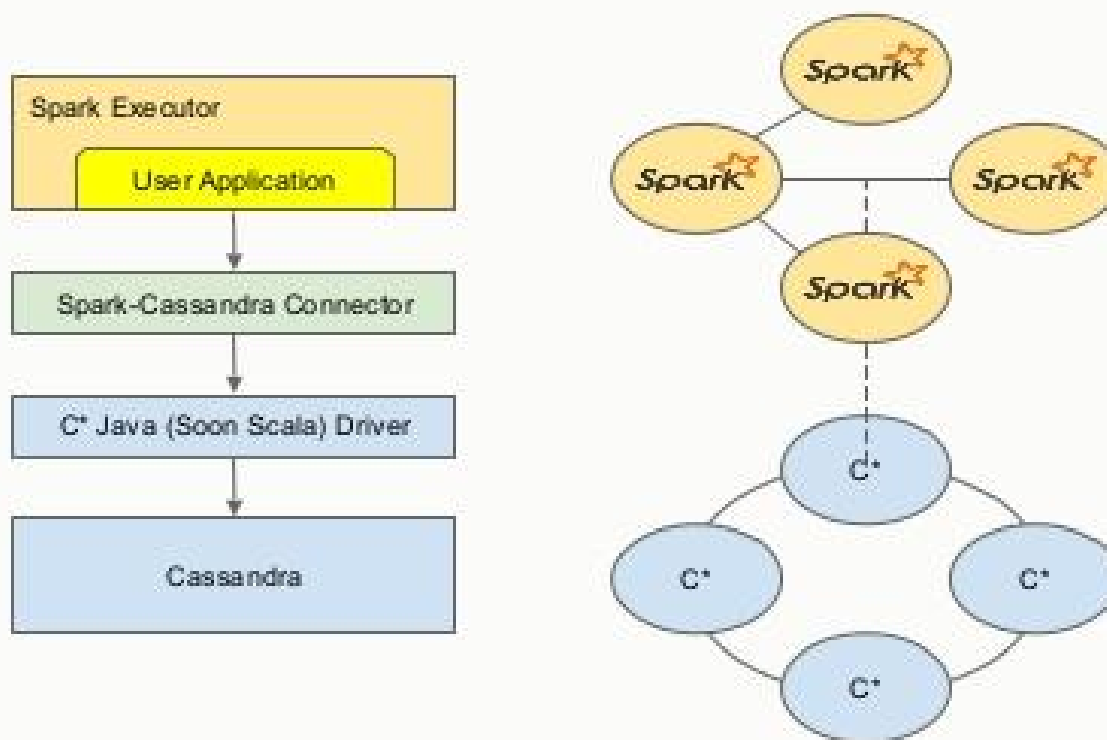
```
val total = sqlContext.sql(  
    "SELECT count(*) FROM subvenciones"  
);  
total.show();
```

```
val group_by_spark = muestra  
    .groupBy("organo_gestor")  
    .sum("importe")  
    .orderBy("sum(importe)")  
group_by_spark.collect.foreach(println)
```


Spark & Cassandra

- Push filters to the server.
- Operations between tables (JOIN, UNION...).
- Know where the data lives (Data Locality).
- Special structures too, like “time-series”.

Spark Cassandra Connector

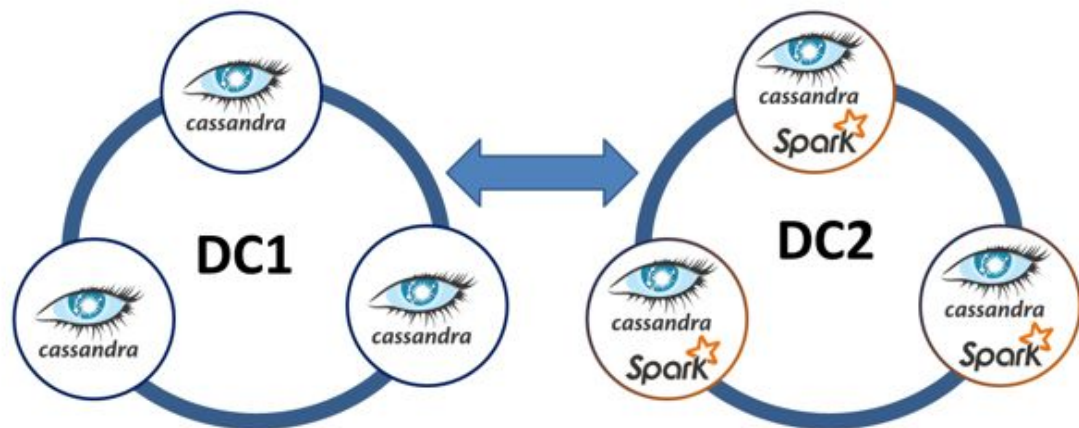


<https://github.com/datastax/spark-cassandra-connector>



Single DC

V.S.



Multi DC

One of the best things - Data Locality

- Any node can coordinate a read/write.
- The connector know where the data lives.
 - Any Spark node reads from his Cassandra node.
- Less use of **CPU** and **RAM** for each node.

I remember nothing!!!

Where can I learn about it?

- Cassandra:

<https://academy.datastax.com/courses/>

- Spark:

<https://databricks.com/training>

<https://community.cloud.databricks.com/>

- Cassandra + Spark:



<https://academy.datastax.com/courses/getting-started-apache-spark>

Use the notebooks!

Importe por meses

```
> %sql
Select sum(importe) AS suma, year(to_date(fecha_concesion)) as year,
month(to_date(fecha_concesion)) as month
from subvenciones
group by year(to_date(fecha_concesion)), month(to_date(fecha_concesion))
Order by year(to_date(fecha_concesion)), month(to_date(fecha_concesion));
```

► (1) Spark Jobs



Spark + Cassandra

- Complementary systems.
- Very scalable and fast.
- Complete Analytical possibilities (Window functions, SQL, ML ...).
- Do you have another database?
 - Do a JOIN with a Cassandra table!

Thanks!

Any question?

You can find me:



Koletzilla



Koletzilla



josemmunoz



Sources

- Slide theme: Escalus
<http://www.slidescarnival.com/es/plantilla-gratis-de-presentacion-escalus/967>
- “The need to lead in data and analytics” :
<http://www.mckinsey.com/business-functions/digital-mckinsey/our-insights/the-need-to-lead-in-data-and-analytics>
- Left image in pag 13: <https://www.youtube.com/watch?v=jxwtmMTMrJA>
- CAP image:
https://www.researchgate.net/figure/282519669_fig1_Figure-1-CAP-theorem-with-databases-that-choose-CA-CP-and-AP
- Netflix Benchmark: <http://techblog.netflix.com/2011/11/benchmarking-cassandra-scalability-on.html>
- Pag 18, extracted from: <http://cassandra.apache.org/>
- Pag 19, benchmarks:
<https://academy.datastax.com/planet-cassandra/NoSQL-performance-benchmarks>

Sources

- Pag 22:
<https://jaceklaskowski.gitbooks.io/mastering-apache-spark/content/spark-architecture.html>
- Pag 24: <https://docs.databricks.com/spark/latest/gentle-introduction/for-data-engineers.html>
- Pag 26: <https://www.datanami.com/2015/11/30/spark-streaming-what-is-it-and-whos-using-it/>
- Pag 30:
<https://www.instaclustr.com/blog/2016/01/07/multi-data-center-apache-spark-and-apache-cassandra-benchmark/>